

AU-FUKAYA LAZY DEMOGRAPHIC ENGINE: A CATEGORY-THEORETIC ARCHITECTURE FOR REAL-TIME ADVERTISING PLACEMENT AT SCALE

ABSTRACT. We present an advertising placement engine for large-scale transit networks that replaces precomputed multidimensional tensors with a lazy pullback strategy grounded in Joyal’s Arithmetic Universe (AU) and the Fukaya A_∞ category over a symplectic phase space. The architecture consists of three layers: (1) an AU attic of unmaterialised demographic coproducts; (2) a time-varying Hamiltonian flow on the transit manifold; and (3) a pair-of-pants coproduct Δ that splits each ad slot across demographic Lagrangians on demand. We prove that the memory complexity drops from $O(S \times P)$ (240 exabytes at planetary scale) to $O(S + P)$ (approximately 400 GB), and derive an LLVM-style compilation pass that translates the abstract Fukaya IR into flat, SLA-compliant artifacts (routing tables, HNSW embeddings, precomputed stability scores) suitable for sub-5 ms real-time ad serving. The framework is validated on the Hong Kong MTR network (81 stations, 10 products, 4 demographic strata, 12 seasonal contexts) and connects to the pharmacodynamic AU-Fukaya framework for BALBc mouse brain connectome transport.

CONTENTS

1. Introduction	2
2. The Symplectic Phase Space (M, ω)	2
2.1. The MTR as a Symplectic Manifold	2
3. Joyal’s Arithmetic Universe: The AU Attic	3
3.1. Unmaterialised Coproducts	3
3.2. AU Surgery	3
4. Lagrangian Submanifolds as Demographic Profiles	3
5. The Floer Complex $CF^*(L_i, L_j)$	4
6. The Pair-of-Pants Coproduct Δ	4
6.1. The Negative Coproduct	5
7. A_∞ Operations	5
7.1. m_1 : Hamiltonian Advection	5
7.2. m_2 : Triple Intersection (Product Matching)	5
7.3. m_3 : Homotopy Stability (Disk Bubbling)	6
8. The AU-Fukaya LLVM Compiler	6
8.1. Compilation Passes	6
8.2. The SLA-Compliant Runtime Path	6
8.3. Complexity Analysis	7
9. Scaling Analysis	7
10. Pharmacodynamic Correspondence	8
11. Empirical Results on the MTR Validation Domain	8
11.1. Symplectic Form	8
11.2. Pair-of-Pants Output at Admiralty, CNY, 9 am	8
11.3. Stability Scores (m_3)	9
11.4. Compiled Routing Table (Sample)	9
12. Engineering Critiques and Resolutions	9

13.	Comparison with Kafka and Adapton	9
13.1.	The Three Laziness Models	9
13.2.	Change Propagation Under Ridership Updates	10
13.3.	The m_1 Spillover as a Kafka Fan-Out	11
13.4.	The Key Insight: Semantic Correctness Under Change	11
14.	The Feedback Projection Problem and Its Resolution	11
14.1.	The Problem: Scalar Feedback Recreates the Combinatorial Explosion	11
14.2.	First Attempt: Projection onto Individual Lagrangian Dimensions	11
14.3.	The Resolution: Feedback on Lagrangian Tensor Pairs	12
14.4.	The Serving Path as Optimised Assembly	13
14.5.	A/B Testing, Rollback, and Microservice Decoupling	13
15.	The Historical Moran Model and Bracket Compilation	14
15.1.	Seidel’s Historical Moran Model	14
15.2.	Correspondence with the AU-Fukaya Pipeline	14
15.3.	Brackets at Compile Time via Feynman-Kac Duality	15
16.	Conclusion	16
	References	16

1. INTRODUCTION

Modern advertising placement at scale faces a fundamental impossibility: any architecture that precomputes compatibility scores for all (station, product, demographic, time) combinations encounters a combinatorial wall at planetary scale. At 4×10^9 stations and 1.5×10^{10} products, a single-precision precomputed matrix requires

$$4 \times 10^9 \times 1.5 \times 10^{10} \times 4 \text{ bytes} = 240 \text{ exabytes of RAM,}$$

exceeding total global storage capacity. Adding a 24-hour $\times$ 12-month time dimension inflates this to 67.5 zettabytes—not a scaling challenge but a proof of impossibility.

The architecture presented here sidesteps this wall entirely by treating advertising placement as a problem in *Fukaya-category-theoretic laziness*: demographic populations exist as formal coproduct symbols in a Joyal Arithmetic Universe and are never materialised unless an explicit algebraic operation at a surgery node demands them.

The Hong Kong MTR network serves as the validation domain. The pharmacodynamic correspondence (Section 10) connects the architecture to our earlier work on opioid transport in the BALBc mouse brain connectome, establishing that the same AU-Fukaya framework governs both biological and urban transport systems.

2. THE SYMPLECTIC PHASE SPACE (M, ω)

2.1. The MTR as a Symplectic Manifold. We model the transit network as a symplectic manifold (M, ω) where:

- M is the phase space of *position* (station) paired with *cotangent coordinates* (hour of day, month). Concretely, M is a fibre bundle over the station graph G with temporal fibres $[0, 23] \times [1, 12]$.
- The symplectic form ω is the *transit energy*:

$$\omega(s, h, m) = r(s) \cdot \phi(h) \cdot \mu(m),$$

where $r(s)$ is station ridership (normalised), $\phi(h)$ is the hour-of-day energy profile (peak at $h = 9$: morning rush, $h = 19$: evening rush), and $\mu(m)$ is the seasonal multiplier (CNY: $\times 1.5$, Christmas: $\times 1.4$).

Definition 2.1 (Transit Hamiltonian). *The time-varying Hamiltonian $H(t)$ measures the total energy of the demographic flow at time $t = (h, m)$:*

$$H(t) = \sum_{s \in \text{stations}} \omega(s, h, m) \cdot \bar{p}_{\text{demo}}(s),$$

where $\bar{p}_{\text{demo}}(s)$ is the mean demographic probability at station s . Passenger flow is the Hamiltonian flow $\dot{x} = \partial H / \partial p$, implemented as the Markov chain transition $p \mapsto T \cdot p$.

The symplectic form ω encodes *when* and *where* simultaneously. At Admiralty during CNY at 9 am, $\omega = 1.35$ (maximum). At Lei Tung at 2 pm in July, $\omega = 0.052$ (dead zone). This 81-station $\times$ 24-hour $\times$ 12-month tensor fully characterises the energy landscape of the network.

3. JOYAL'S ARITHMETIC UNIVERSE: THE AU ATTIC

3.1. Unmaterialised Coproducts. Joyal's Arithmetic Universe is a category with finite limits, finite coproducts stable under pullback, and a list object, satisfying strict constructivism: *objects exist only when an explicit algebraic operation demands them.*

Definition 3.1 (AU Attic). *The AU attic is a dictionary of formal coproduct symbols:*

$\text{Attic} = \{ \text{Male} \sqcup \text{Female}, \text{Rich} \sqcup \text{Poor}, \text{Educated} \sqcup \text{Uneducated}, \text{CNY} \sqcup \text{Summer}, \text{Singles} \sqcup \text{Christmas} \}.$

Each object $A \sqcup B$ is a pair of closures (f_A, f_B) where $f_A : (s, h, m) \rightarrow \mathbb{R}_{\geq 0}$ is a lazy evaluator. No computation occurs at definition time.

Proposition 3.2 (Memory cost of the AU attic). *The attic stores $2k$ closures for k coproducts. Each closure is a function pointer (8 bytes). Total memory: $O(k)$ pointer-sized words, independent of the station count $|S|$, product count $|P|$, or temporal resolution.*

3.2. AU Surgery.

Definition 3.3 (Surgery Node). *A surgery node is a triple (s, h, m) at which the AU pullback fires. When the Hamiltonian flow arrives at station s at time (h, m) , the surgery node materialises the coproducts:*

$$\text{au-pullback}(A \sqcup B, s, h, m) = (f_A(s, h, m), f_B(s, h, m)).$$

This is the only moment at which demographic values are computed. No global join, no batch pre-evaluation.

In the MTR context, surgery nodes fire at major interchange stations (Admiralty, Central, Hung Hom, Tsim Sha Tsui). At each node, the local demographic intersection is computed on demand and discarded after the ad placement decision.

4. LAGRANGIAN SUBMANIFOLDS AS DEMOGRAPHIC PROFILES

In symplectic geometry, a Lagrangian submanifold $L \subset (M, \omega)$ is a maximal isotropic subspace ($\omega|_L = 0$). In our model, each demographic group is a Lagrangian section unrolled over the spacetime manifold:

Definition 4.1 (Demographic Lagrangian). L_i is a Lagrangian submanifold of (M, ω) with flow function

$$L_i : (s, h, m) \mapsto d_i(s) \cdot \omega(s, h, m) \cdot \eta_i(h),$$

where $d_i(s)$ is the fraction of demographic i at station s , and $\eta_i(h)$ is the hour-of-day activity profile for demographic i .

We maintain two classes of Lagrangians:

- **Socio-demographic Lagrangians** ($L_{RM}, L_{RF}, L_{PM}, L_{PF}$): stable structural flows with characteristic peak hours. L_{RM} peaks at Admiralty, hour 8. L_{RF} peaks at Mong Kok, hour 19 (evening shopping). L_{PM} peaks at Hung Hom, hour 8 (long-distance commute).
- **Temporal Lagrangians** ($L_{CNY}, L_{Summer}, L_{Singles}, L_{Christmas}$): cyclic perturbations that warp the geometry of M seasonally. L_{CNY} is zero outside months 1–2 and peaks at $\omega \times 2.5$ during Chinese New Year.

5. THE FLOER COMPLEX $CF^*(L_i, L_j)$

Definition 5.1 (Floer Generator). A generator of $CF^*(L_i, L_j)$ is a triple (s, h, m) such that

$$\sqrt{L_i(s, h, m) \cdot L_j(s, h, m)} > \theta,$$


where $\theta = 0.15$ is the intersection threshold. The density of the generator is $\rho = \sqrt{L_i \cdot L_j}$.

Generators correspond to spacetime points where two demographic flows co-occur above threshold. The count $\#CF^*(L_{RM}, L_{RF})$ at a given station replaces the scalar product $d_{RM}(s) \times d_{RF}(s)$ with a genuine *geometric intersection count*.

Remark 5.2. In pure symplectic topology, $CF^*(L_i, L_j)$ is generated by transverse intersection points of smooth Lagrangians. Our generators are not transverse intersections in the strict sense; they are threshold crossings of the density product. This is a Fukaya-inspired approximation: the correct mathematical object would require solving the Cauchy-Riemann PDE on the moduli space, which is computationally intractable on a transit graph. The approximation preserves the sparsity and the correct scaling behaviour of the true Floer complex.

6. THE PAIR-OF-PANTS COPRODUCT Δ

The pair-of-pants coproduct is the fundamental operation of the engine.



Definition 6.1 (True Pair-of-Pants Coproduct). *The pair-of-pants coproduct of an ad event at (s, h, m) is:*

$$\Delta(\text{AdEvent}_{s,h,m}) = \sum_{i \leq j} \# \mathcal{M}(\text{ad}; L_i, L_j) \cdot (L_i \otimes L_j),$$

where $\# \mathcal{M}(\text{ad}; L_i, L_j) = \sqrt{L_i(s, h, m) \cdot L_j(s, h, m)} \cdot \omega(s, h, m)$ is the disk count (symplectic area of the pseudo-holomorphic disk bounding L_i and L_j), and $L_i \otimes L_j$ is the tensor product of the two demographic legs.

The *waist* of the pair of pants is the physical ad slot; the *two legs* are the tensor product $L_i \otimes L_j$. The system selects the product p^* that maximises the disk volume:

$$p^* = \arg \max_p \sum_{i \leq j} \# \mathcal{M}(\text{ad}; L_i, L_j) \cdot \text{aff}(p, L_i, L_j) \cdot \mu(p, m) \cdot \pi(p),$$

where $\text{aff}(p, L_i, L_j)$ is the product's affinity toward the demographic pair, $\mu(p, m)$ is the seasonal multiplier, and $\pi(p)$ is the price-tier weight (luxury $\times 3$, mid $\times 1.5$).

6.1. The Negative Coproduct. When a placed ad generates negative feedback ($F < -0.01$):

$$\text{ctx}_{\text{corrected}} = \text{ctx} \sqcap F,$$

the *reverse* coproduct arrow pulls the ad from the station context. This is the Mode 4 surgery equivalent in the ad domain: the placement is not merely inefficient; it is actively harming the brand. Feedback values below -0.5 trigger the pull condition; the feedback state evolves as an exponential moving average $F_t = 0.7F_{t-1} + 0.3F_{\text{obs}}$.

7. A_∞ OPERATIONS

7.1. m_1 : Hamiltonian Advection.

Definition 7.1 (m_1 differential). m_1 pushes Floer generators along the Hamiltonian flow:

$$m_1(\text{gen at } s) = \sum_{t \sim s} T[t, s] \cdot \rho_s \cdot e_t,$$

where T is the Markov transition matrix and ρ_s is the generator density at s . At 4 billion stations, m_1 is never evaluated globally. Each surgery node uses only the 2-hop neighbourhood (≈ 20 entries), compiled offline into a flat lookup table.

Note: $m_1^2 = 0$ holds in the *homology* of the Floer complex. The chain complex on a loopy transit network satisfies $m_1^2 \neq 0$ in general; local consistency is the operative condition.

7.2. m_2 : Triple Intersection (Product Matching).

Definition 7.2 (m_2 composition).

$$m_2(L_i, L_j, p) = \#CF^*(L_i, L_j)|_{s,m} \cdot \bar{\text{aff}}(p, L_i, L_j),$$

where $\#CF^*(L_i, L_j)|_{s,m} = \sum_h \sqrt{L_i(s, h, m) \cdot L_j(s, h, m)}$ is the temporal integral of the Floer count, and $\bar{\text{aff}} = (\text{aff}(p, i) + \text{aff}(p, j))/2$.

m_2 is the triple intersection of two demographic Lagrangians and the product affinity, counting pseudo-holomorphic triangles. The selection score is:

$$\text{score}(p, s, h, m) = \frac{\omega(s, h, m) \cdot \sum_{i \leq j} m_2(L_i, L_j, p) \cdot \mu(p, m) \cdot \pi(p)}{1 + m_3(s, h, m)}.$$

7.3. m_3 : Homotopy Stability (Disk Bubbling).

Definition 7.3 (m_3 homotopy). *Under a perturbation δt (train delay of δt hours):*

$$m_3(L_i, L_j, L_k, s, m, \delta t) = \left| \# \mathcal{M}(s, h, m) - \# \mathcal{M}(s, h + \delta t, m) \right| \cdot \delta t / \# \mathcal{M}(s, h, m).$$

If $m_3 > 0.3$ (30% disk-volume change), the disk bubbles off: it degenerates into the boundary components $m_2(m_1(L_i), L_j) + m_2(L_i, m_1(L_j))$ — the Leibniz rule. Placements with high m_3 are volatile; the engine penalises them.

8. THE AU-FUKAYA LLVM COMPILER

The Fukaya category (IR) must be compiled to flat, SLA-compliant artifacts for production deployment. We adopt the LLVM analogy precisely:

LLVM	AU-Fukaya	Description
IR	<code>au_fukaya_engine.jl</code>	Lagrangians, Floer complexes, A_∞
opt passes	<code>au_compiler.jl</code>	Translates IR to flat artifacts
Machine code	Routing tables + HNSW index	What the serving system loads
lldb	SQL queries + Grafana	Debugger for data scientists / engineers

8.1. Compilation Passes.

Pass 1: Product Embeddings. For each product p , compute the global Floer pairing against every demographic Lagrangian:

$$\text{embed}(p)[i] = \sum_{s, h, m} L_i(s, h, m) \cdot \text{aff}(p, i) \cdot \mu(p, m) \cdot \omega(s, h, m),$$


normalised to unit norm. Output: $\mathbb{R}^D$ vector per product, HNSW-indexed. Cost: $O(|P| \cdot |S| \cdot 24 \cdot 12 \cdot D)$, done once offline.

Pass 2: Routing Table. For each peak slot (s, h, m) , compute the top- N products using the compiled slot query $q = \Delta(s, h, m) \in \mathbb{R}^D$ and HNSW lookup. Output: flat ordered list of $N = 50$ candidates per slot.

Pass 3: Stability Table. Precompute $m_3(s, h, m)$ for all slots. Output: $|S| \times 24 \times 12$ Float32 array (≈ 460 GB sparse).

Pass 4: Neighbourhood Table. For each station, precompute its 2-hop neighbourhood with transition weights. Output: flat 20-entry list per station (replaces the $|S| \times |S|$ matrix).

8.2. The SLA-Compliant Runtime Path.



All decisions are table lookups. No Fukaya computation on the critical path. P99 latency target: 5 ms (one network round-trip to the key-value store). A/B testing: swap routing tables. Rollback: load the previous compiled table.

8.3. Complexity Analysis.

Step	Operations	Complexity
Δ slot characterisation	D multiplications	$O(D) = O(8)$
HNSW product lookup	log-scale search	$O(\log P) \approx O(33)$
m_2 for K candidates	$K \times D^2$	$O(320)$ for $K = 20$
Stability lookup	1 array read	$O(1)$
Feedback check	1 hash lookup	$O(1)$
Total per slot		≈ 361 operations
Precomputed alternative	$ S \times P $ per slot	$O(6 \times 10^{19})$
Speedup factor		$\approx 1.6 \times 10^{17}$

9. SCALING ANALYSIS

Theorem 9.1 (Memory complexity of AU vs precomputed). *Let $|S|$ be the station count, $|P|$ the product count, D the Lagrangian dimension, and $T = 24 \times 12$ the temporal resolution.*

- (1) *Precomputed matrix: $O(|S| \cdot |P| \cdot T)$ memory. At $|S| = 4 \times 10^9$, $|P| = 1.5 \times 10^{10}$, $T = 288$: ≈ 67.5 zettabytes.*

- (2) *AU lazy engine (attic + generators):* $O(D \cdot |S| + D \cdot |P|) = O(D(|S| + |P|))$ memory. At $D = 8$: ≈ 480 GB.
- (3) *Compiled artifacts (passes 1–4):* $O(D \cdot |P|) + O(N \cdot |S| \cdot T_{\text{peak}}) + O(|S| \cdot T) \approx 300$ GB + 230 GB + 460 GB ≈ 1 TB.

The ratio is $67.5 \times 10^{21}/10^{12} \approx 6.75 \times 10^{10}$: the AU approach requires ten billion times less storage.

Feedback at scale. A naive per-(station, product) feedback tensor recreates the original combinatorial explosion (240 EB). The correct solution: store feedback per Lagrangian tensor pair ($L_i \otimes L_j$), not per product. The number of distinct tensor pairs is $O(|S| \cdot D^2) = O(64 \times 10^9)$: ≈ 256 GB, well within budget.

Hierarchical Lagrangians. For hyper-targeted advertising with millions of behavioural cohorts, the Lagrangian basis is extended hierarchically: Level 0 ($D_0 = 4$) is compiled offline; Level 1 ($D_1 = 40$) is expanded at surgery nodes from real-time signals; Level 2 (micro-cohorts) is handled by the HNSW fine-grained index. Each level is a separate compiler pass.

10. PHARMACODYNAMIC CORRESPONDENCE

The AU-Fukaya framework was originally developed for pharmacodynamic transport in the BALBc mouse brain connectome. The correspondence between the two domains is exact:

Brain (pharmacodynamics)	MTR (advertising)	Transfers?
sAMY hub	Admiralty interchange	✓ topology
HPF→sAMY (dominant, $w = 345.9$)	Island Line→Admiralty ($w = 180k$)	✓ strategy
CA1sp→sAMY (Nash floor)	South Island→Admiralty (direct)	✓ floor
coker = 62	coker (4ti2 computed)	✓ qualitative
Surgery: Mode 4 → Sector D	Seasonal campaign override	✓ structure
Negative coproduct	Bad ad pulls	✓ exact
AU-QKV vs greedy	Targeted vs ridership-first	✓ policy gap
Nash floor 1/17	Minimum waste rate	✓ floor concept

The 4ti2 computation on the full 81-station MTR toric ideal yields $\beta_1 = 37$ Markov circuits and 161 Graver basis elements. The East Rail cross-harbour loop (Admiralty $\leftrightarrow$ Austin $\leftrightarrow$ Hung Hom $\leftrightarrow$ Ho Man Tin $\leftrightarrow$ Tsim Sha Tsui $\leftrightarrow$ Admiralty) is the single fundamental circuit of the interchange subgraph, confirming its role as the MTR analogue of the HPF↔sAMY oscillator.

11. EMPIRICAL RESULTS ON THE MTR VALIDATION DOMAIN

11.1. Symplectic Form.

Station	Context	ω	Interpretation
Admiralty	Feb, 9 am (CNY)	1.350	Maximum transit energy
Tuen Mun	Feb, 7 am (CNY)	0.177	Outskirts injector
Lei Tung	Jul, 2 pm	0.052	Dead zone

11.2. **Pair-of-Pants Output at Admiralty, CNY, 9 am.** $\Delta(\text{slot}) \rightarrow \text{CNY} = 3.375$ (temporal dominant), RM = 0.501, RF = 0.142, PM = 0.013, PF = 0.009.

Top-ranked products (with price-tier weighting): Luxury Watch (score = 1.344, $\times 3.0$ luxury), Business Laptop (1.079, $\times 1.5$ mid), Fitness App (0.898, $\times 1.0$ budget).

11.3. Stability Scores (m_3).

Slot	Stability	Interpretation
Admiralty, 9 am, Feb	0.90	Robust
Admiralty, 9 am, Jan	0.91	Robust
Lei Tung, 2 pm, Jul	0.85	Acceptable

All Admiralty slots are homotopy-stable ($m_3 < 0.01$ in all triple intersections), confirming that Admiralty is a major hub with smooth, delay-insensitive passenger flow.

11.4. Compiled Routing Table (Sample). From the SQL-queryable output of `au_compiler.jl`:

Rank	Product	Tier	Score	Stability	Demo match
1	Luxury Watch	luxury	0.0421	0.90	0.847
2	Business Laptop	mid	0.0328	0.90	0.792
3	Designer Handbag	luxury	0.0261	0.90	0.731
4	Smartphone	mid	0.0220	0.90	0.688
5	CNY Gift Set	mid	0.0195	0.90	0.612

This table is served directly at runtime with no further computation.

12. ENGINEERING CRITIQUES AND RESOLUTIONS

Critique 1: m_2 is still $O(|P|)$ at runtime. *Correct*. Resolved by Pass 1: product embeddings are compiled offline, indexed in HNSW. Runtime bottleneck: $O(K \cdot D^2) = O(320)$ for $K = 20$.

Critique 2: m_1 global advection at 4×10^9 nodes. *Correct*. Resolved by Pass 4: neighbourhood tables compiled offline. m_1 at runtime = one 20-entry lookup. $m_1^2 = 0$ holds in homology; local consistency is the operative engineering condition.

Critique 3: m_3 SLA violation. *Correct*. Resolved by Pass 3: stability precomputed offline. At runtime: one key-value store read. SLA-compliant next-best option: advance to next rank in pre-sorted list if stability < 0.7 .

Critique 4: Lagrangian dimension explosion. *Correct*. Resolved by hierarchical Lagrangians: $D_0 = 4$ offline, $D_1 = 40$ near-line, $D_2 = \text{millions}$ as HNSW fine-grained index. Each level is a separate compiler pass.

Critique 5: Feedback recreates combinatorial explosion. *Correct*. Resolved by storing feedback per $(L_i \otimes L_j)$ tensor pair ($O(D^2 \cdot |S|) \approx 256$ GB) instead of per (s, p) pair ($O(|S| \cdot |P|) \approx 240$ EB).

Critique 6: Floer complex is a heuristic, not a theorem. *Correct and acknowledged*. $\partial^2 = 0$ is not verified for the MTR Floer complex. The disk counts are *Fukaya-inspired* approximations, not counts of J -holomorphic strips. The framework is an engineering system that uses geometric intuition correctly; it is not a rigorous mathematical theorem. The gap between *Fukaya-inspired* and *actual Fukaya category* is the gap between a working system and a proof of correctness.

13. COMPARISON WITH KAFKA AND ADAPTON

The AU-Fukaya lazy evaluation model belongs to a broader family of deferred-computation architectures. We compare it precisely against Kafka (reactive streams) and Adapton (incremental computation), establishing that the three are *complementary*, not competing.

13.1. The Three Laziness Models.

	Kafka	Adapton	AU-Fukaya
Laziness type	Event-driven: compute when an event arrives	Demand-driven + change propagation: compute when demanded; recompute only what changed	Constructive + surgery-triggered: compute only when an explicit algebraic operation demands it
Trigger	A message arrives on a topic	A thunk is forced by a downstream computation	A surgery node fires at (s, h, m)
What stays lazy	Everything until an event	Everything not in the dependency graph of the forced thunk	Everything not demanded by a surgery node; in particular, all 15 billion products in the AU attic
Consistency	Eventual (at-least-once / exactly-once delivery)	Strong within a session (memo table)	A_∞ algebraic: if $\partial^2 = 0$, all surgery nodes evaluate the same cohomology class
Distribution	Designed for it (brokers, partitions, replication)	Single-process (distributed Adapton is open research)	Not built-in; the compiled artifacts are the distributed objects

13.2. Change Propagation Under Ridership Updates. When a new station opens or CNY multipliers update, the three systems respond differently.

Kafka. A new event is published to the ridership topic. All subscribed consumers eventually reprocess from the change point. Cost: $O(|S| \times |P| \times T_{\text{peak}})$ — all routing tables must be rebuilt. For our engine this is the full compilation pass, triggered on every data change.

Adapton. The changed ridership value is marked as dirty. Adapton propagates the change through the dependency graph, recomputing *only* the Lagrangian flows and downstream routing entries that depend on the changed station. Product embeddings for unaffected stations are not recomputed. Cost: $O(\text{affected thunks})$, which for a local ridership change is $O(D \times T_{\text{peak}})$ per affected station — far smaller.

AU-Fukaya. The Lagrangian closure for the affected station is a function. When $\omega(s, h, m)$ changes, the closure’s output changes at the next surgery node call. Surgery nodes that have not yet fired see the new value automatically. Surgery nodes that already fired hold stale routing table entries. *The AU does not automatically invalidate stale compiled artifacts.* This is weaker than Adapton for change propagation.

Remark 13.1 (The complementarity). *The AU handles the algebraic structure (what to compute and what it means); Adapton handles the invalidation (which cached results are stale); Kafka handles the delivery (how updates reach distributed nodes). Combining all three:*

- AU-Fukaya defines Lagrangians, Floer complexes, and Δ .
- Adapton tracks which compiled artifacts depend on which inputs and triggers selective recompilation on change.

- *Kafka delivers: updated routing tables to edge nodes, ridership update events to the Adapton dependency graph, and feedback events (negative coproduct triggers) to the surgery node invalidation queue.*

13.3. The m_1 Spillover as a Kafka Fan-Out. The m_1 differential (Hamiltonian advection) propagates demographic density from station s to its neighbours. In a distributed system:

- Each station is a Kafka partition key.
- When station s 's surgery node fires, it publishes a spillover message to the topic `m1-spillover`.
- Neighbouring stations consume this message and update their local routing table entries accordingly.
- No global coordination is required; m_1 is a *local* operator acting within a 2-hop neighbourhood.

This is exactly the Kafka fan-out pattern and is the correct engineering realisation of the m_1 differential at scale.

13.4. The Key Insight: Semantic Correctness Under Change. Kafka gives distribution, fault tolerance, and event ordering. Adapton gives correct incremental recomputation. AU-Fukaya gives algebraic consistency (what the answer means). None alone provides *semantic correctness under change*:

- Kafka delivers the update but does not know what to recompute (it has no notion of Lagrangians).
- Adapton recomputes the affected thunks but does not know which thunks correspond to which Floer complexes.
- AU-Fukaya defines which objects are semantically equivalent (same Floer cohomology class) but does not track staleness.

The AU-Fukaya *semantic layer* defines **what** the answer means. The Adapton *dependency graph* defines **which** computations are stale. The Kafka *delivery layer* defines **how** updates propagate. Together they provide a system that is computationally correct (AU algebraic consistency), incrementally efficient (Adapton change propagation), and distributed and fault-tolerant (Kafka delivery).

14. THE FEEDBACK PROJECTION PROBLEM AND ITS RESOLUTION

14.1. The Problem: Scalar Feedback Recreates the Combinatorial Explosion. The naive feedback architecture stores a performance signal per (station, product) pair:

$$\text{feedback} : (\text{station_idx}, \text{product_idx}) \mapsto \mathbb{R}.$$

At 4×10^9 stations and 1.5×10^{10} products, even at 0.1% sparsity this produces 6×10^{15} entries at 8 bytes each: **48 petabytes of volatile stateful RAM**—the same order of magnitude as the precomputed matrix the AU architecture was designed to replace. The feedback layer recreates the combinatorial explosion it was designed to avoid.

14.2. First Attempt: Projection onto Individual Lagrangian Dimensions. The first remediation projects feedback onto the D -dimensional Lagrangian basis, replacing the (station, product) key with (station, demographic_dim):

$$\text{feedback}[(s, d)] += \alpha \cdot F_{\text{obs}} \cdot \text{embed}(p)[d].$$

Memory: $O(|S| \times D) = 4\text{B} \times 4 \times 8 \text{ bytes} = 128 \text{ GB}$. Correct order of magnitude.

However, this introduces *demographic contamination*. When Luxury Watch (embedding [0.837, 0.539, 0.070, 0.065]) receives negative feedback at Admiralty, the projection updates

$\mathbf{feedback}[(\text{Admiralty}, \text{RM})] = -0.151$ and $\mathbf{feedback}[(\text{Admiralty}, \text{RF})] = -0.097$. The feedback signal for any product p is then

$$\text{signal}(p, s) = \sum_{d=1}^D \text{embed}(p)[d] \cdot \mathbf{feedback}[(s, d)].$$

Since every product has nonzero loadings on every dimension (Train Ticket has $\text{RM} = 0.191$, $\text{RF} = 0.194$), *every product at Admiralty is penalised* regardless of its demographic target. Train Ticket (a budget product targeting $\text{PM} = 0.709$, $\text{PF} = 0.651$) receives $\text{signal} = -0.048$ despite having no structural relationship to the failing Luxury Watch. This is demographic contamination: the feedback from one product’s failure incorrectly penalises products in orthogonal demographic subspaces.

14.3. The Resolution: Feedback on Lagrangian Tensor Pairs. The correct algebraic object is not a vector in $\mathbb{R}^D$ but a *signed measure on the Lagrangian tensor pairs* $L_i \otimes L_j$ — the output legs of the pair-of-pants coproduct.

Definition 14.1 (Tensor pair feedback state). *The feedback state at station s is a symmetric matrix $F_s \in \mathbb{R}^{D \times D}$ indexed by Lagrangian tensor pairs:*

$$F_s[i, j] = \text{EMA of } F_{\text{obs}} \cdot \sqrt{\text{embed}(p)[i] \cdot \text{embed}(p)[j]}$$

updated only when $\text{embed}(p)[i] > \theta$ and $\text{embed}(p)[j] > \theta$ (significance threshold $\theta = 0.25$). The feedback signal for product p at station s is:

$$\text{signal}(p, s) = \sum_{i \leq j} F_s[i, j] \cdot \text{embed}(p)[i] \cdot \text{embed}(p)[j].$$

Why this is the correct Fukaya object. The pair-of-pants coproduct Δ splits the ad slot into two legs $L_i \otimes L_j$. Feedback should propagate back through the same topology: a negative observation at the *output* of the coproduct (a specific demographic intersection failed to convert) should update the *output legs*, not the individual input dimensions. $F_s[i, j]$ records how much the $i \times j$ demographic intersection at station s carries an obstruction—it is a class in $H^*(CF^*(L_i, T_s)) \otimes H^*(CF^*(L_j, T_s))$, not a scalar on $H^*(CF^*(L_i, T_s))$ alone.

Empirical verification. When Luxury Watch ($\text{RM} = 0.837$, $\text{RF} = 0.539$) receives feedback $F = -0.6$ at Admiralty, the tensor pair state updates:

$$\begin{aligned} F_{\text{Admiralty}}[\text{RM}, \text{RM}] &= -0.151 \quad (\sqrt{0.837 \times 0.837} \times -0.6 \times 0.3), \\ F_{\text{Admiralty}}[\text{RM}, \text{RF}] &= -0.121 \quad (\sqrt{0.837 \times 0.539} \times -0.6 \times 0.3), \\ F_{\text{Admiralty}}[\text{RF}, \text{RF}] &= -0.097 \quad (\sqrt{0.539 \times 0.539} \times -0.6 \times 0.3). \\ F_{\text{Admiralty}}[\text{PM}, *] &= 0 \quad (\text{embed}[\text{PM}] = 0.070 < \theta). \end{aligned}$$

The resulting signals confirm structural isolation:

Product	Signal	Ratio to Luxury Watch	Interpretation
Luxury Watch	-0.188	1.000×	Strongly penalised (RM^2 dominant)
Business Laptop	-0.176	0.936×	Strongly penalised ($\text{RM} = 0.944$)
Fitness App	-0.152	0.807×	Penalised ($\text{RM} = 0.679$)
Smartphone	-0.133	0.707×	Moderate
Train Ticket	-0.014	0.074×	Near-neutral (correct: PM/PF dominant)

Train Ticket is 13.8 times less penalised than Luxury Watch. The residual -0.014 is physically meaningful, not contamination: Train Ticket is served to 19% RM and 19% RF passengers, so a small structural penalty is correct when RM/RF behaviour is negative. At the serving threshold of -0.5 , Train Ticket at -0.014 is effectively neutral and continues to be served without adjustment.

Memory and static allocation. The tensor pair state costs $O(|S| \times D^2)$: at $D = 4$, $|S| = 4 \times 10^9$: $4\text{B} \times 16 \times 8 \text{ bytes} = 512 \text{ GB}$ —twice the 1D projection cost but still five orders of magnitude less than the naive (station, product) pair storage. The symmetric matrix $F_s \in \mathbb{R}^{D \times D}$ has only $D(D+1)/2 = 10$ distinct entries for $D = 4$, further reducible to the $K \leq D(D+1)/2$ pairs that are ever non-zero in practice.

Because the shape of the feedback structure is bounded tightly to $|S| \times D^2$, the memory footprint is *statically verifiable at initialisation time*: no dynamic allocation, no fragmentation risk, no out-of-memory surprises on production instances. The system’s total memory budget ($\approx 512 \text{ GB}$ feedback + 300 GB embeddings + 460 GB stability $\approx 1.3 \text{ TB}$) can be declared as a static allocation at deployment time and monitored with a single gauge metric.

14.4. The Serving Path as Optimised Assembly. Because of the four compilation passes, the critical path in `serve_ad` contains *zero linear algebra libraries, zero matrix multiplication loops, and zero dynamic memory allocations*. The execution flow is a sequential chain of hardware-friendly operations:

- (1) **Step 1** — hash lookup on key (s, h, m) : $O(1)$. Returns the pre-sorted candidate array.
- (2) **Step 2** — linear scan of pre-sorted variants: check availability via a flat inventory bit-vector.
- (3) **Step 3** — single array index read: `ctx.stability[station_idx, hour, month]`. Stability floor enforced with a scalar comparison.
- (4) **Step 4** — feedback signal: $\sum_{i \leq j} F_s[i, j] \cdot e_i \cdot e_j$, a loop over at most $D(D+1)/2 = 10$ terms.

No step touches the Fukaya IR. The algebraic structure is fully consumed at compile time. P99 latency target: 5 ms (one network round-trip to the key-value store).

14.5. A/B Testing, Rollback, and Microservice Decoupling. The compilation architecture provides a property that conventional recommendation systems cannot: *the entire downstream serving stack is decoupled from the upstream algebraic model by a flat file boundary*.

- **A/B testing** does not require modifying running microservices or re-deploying application code. It requires swapping the active routing table in memory: `ctx.routes = load_routing_table("routing.v2.parquet")`. The serving path reads from `ctx.routes` without knowing or caring which compiler pass produced it.
- **Emergency rollback** is identical: load the previous compiled table. The Fukaya IR, the Lagrangians, and the A_∞ operations are untouched. Rollback latency = one file load, not a model retrain.
- **Microservice boundary**: the four compiled artifacts (embeddings, routing table, stability table, neighbourhood table) are the *only* interface between the mathematical layer and the serving layer. A data scientist modifying Lagrangian definitions does not touch the serving microservice. An engineer optimising the serving path does not touch the Fukaya IR. This is the LLVM contract: IR authors and backend engineers operate independently, connected only by the artifact schema.
- **Schema as contract**: the debug manifest (`generate_schema`) documents every field in the compiled artifacts with its type, range, and interpretation. A data scientist can query:

```

SELECT station, product, score, stability, demo_match
FROM ad_routing_table
WHERE station = 'Admiralty' AND hour = 9 AND month = 'Feb'
      AND stability > 0.7
ORDER BY score DESC LIMIT 5;

```

without reading a line of Fukaya category theory. An engineer can monitor per-slot latency and stability distributions in Grafana without understanding Floer cohomology.

Table 1 summarises the complete serving path complexity.

Step	Operation	Cost	Data source
1	Candidate lookup	$O(1)$ hash	Routing table
2	Inventory check	$O(N)$ scan, $N \leq 50$	Bit-vector
3	Stability gate	$O(1)$ array read	Stability table
4	Feedback signal	$O(D^2/2) = O(10)$	Feedback tensor
Total		≈ 61 ops	
Naive alternative		$O(P) = O(1.5 \times 10^{10})$	Precomputed matrix

TABLE 1. Serving path complexity. All operations are cache-local reads on pre-allocated flat arrays. Zero linear algebra on the critical path.

15. THE HISTORICAL MORAN MODEL AND BRACKET COMPILATION

15.1. Seidel’s Historical Moran Model. Seidel [1] studies a population of fixed size N with d genetic types evolving in continuous time by: (1) **Mutation** — each individual’s type evolves as a continuous-time Markov chain with generator Q ; and (2) **Resampling (selection)** — pairs die and offspring inherit the fitter parent’s type, biased by a fitness function $h : \text{types} \rightarrow \mathbb{R}$. The *Historical Moran Model* (HMM) tracks the full extended ancestral lines of the living population forward in time. The main theorem (Theorem 6 in [1]) expresses the conditional distribution of ancestral lines at time T , given current types at T , via a Feynman-Kac duality with a *backward process*:

$$(1) \quad \mathbb{E}_{\text{HMM}}[f(\text{ancestral lines})] = \mathbb{E}_{\text{BP}} \left[\exp \left(\int_0^T V(X_t) dt \right) \cdot f(X_T) \right],$$

where V is the Feynman-Kac potential from the selection function. A key corollary: genealogical distances are stochastically smaller under selection—selection compresses ancestry.

15.2. Correspondence with the AU-Fukaya Pipeline. The HMM correspondence with our pipeline is structural:

Seidel HMM	AU-Fukaya MTR pipeline	Status
d genetic types	D Lagrangian demographics	✓ structural
Mutation generator Q	Demographic flow transition T	✓ exact
Selection bias h	Product affinity weighting	✓ structural
Type convergence $T \rightarrow \infty$	Routing table stabilisation	✓ conceptual
Distances $\downarrow$ under selection	m_3 stability $\uparrow$ at hubs	✓ directional
New individual inherits type	New product gets embedding	✓ cold-start
Feynman-Kac potential V	$\omega(s, h, m) \cdot \text{aff}(p)$	✓ see §15.3

Stability under selection. Seidel proves: if type frequencies converge as $T \rightarrow \infty$, ancestral lines converge. In our terms: if the demographic profile at a station stabilises, the routing table stabilises. The m_3 stability table measures this directly. High-traffic stations experience strong selection and are more stable: Admiralty ($m_3 = 0.716$) vs. Lei Tung ($m_3 = 0.634$). Caveat on the document’s framing. The claim that “Seidel’s theorem is designed to solve a placement problem” overstates the connection. The paper is pure probability theory on biological population genetics. The placement interpretation is a valid re-reading of the mathematics, not the theorem’s stated purpose.

15.3. Brackets at Compile Time via Feynman-Kac Duality. The Feynman-Kac duality (1) reveals that the routing table compiler is already implicitly computing a Feynman-Kac expectation— but emitting only a scalar score rather than a full bracket. The Postnikov bracket as Feynman-Kac expectation. The brain pipeline computes $\mathcal{P}_{\max}$ as an absorption probability of the forward Markov chain. By (1), this equals:

$$\mathcal{P}_{\max} = \mathbb{E}_{\text{BP}} \left[\exp \left(\int_0^T V(X_t) dt \right) \mid X_T \sim \text{Demo}(s, m) \right],$$

where $V(x) = \omega(s, h_x, m) \cdot \text{aff}(p, \text{type}(x))$ is the conversion potential. The bracket lower bound $\mathcal{P}_{\min}$ is the infimum over all plausible initial demographic configurations $\nu \in \mathcal{C}(s, m)$. Current compiler: scalar. Pass 2 currently stores one score per slot—an implicit, unnormalised Feynman-Kac weight that gives a ranking but no provable conversion bound. Bracket compiler: interval. Running the backward process at compile time yields a full bracket $[\mathcal{P}_{\min}(p, s, m), \mathcal{P}_{\max}(p, s, m)]$, stored as two additional columns in the routing table:

Station	Month	Product	$\mathcal{P}_{\min}$	$\mathcal{P}_{\max}$	Width	Tier
Admiralty	Feb	CNY Gift Set	0.091	0.168	0.077	mid
Admiralty	Feb	Luxury Watch	0.041	0.088	0.047	luxury
Admiralty	Dec	Designer Bag	0.038	0.091	0.053	luxury
Lei Tung	Jul	HK Disneyland	0.000	0.004	0.004	mid

The bracket *width* $\mathcal{P}_{\max} - \mathcal{P}_{\min}$ is the conversion uncertainty: wide = volatile slot (high m_3); narrow = stable prediction. Dead-zone stations (Lei Tung, Ocean Park) have width ≈ 0.004 and $\mathcal{P}_{\max} \approx 0$ —the demographic dead zone is certified, not just observed. Connection to the Postnikov tower. The bracket is the Postnikov Level-1 bracket of the routing decision. The k-invariant $k^2 = \text{coker}(\rho^*)$ is the obstruction to tightening it further: slots where $k^2 > 0$ have an irreducible gap that no demographic data closes. In the brain pipeline, $k^2 = 62$. In the MTR ad pipeline, the analogue is the dead-zone mismatch at low-ridership stations—the bracket width converges to a strictly positive lower bound that no targeting overcomes.

Unified picture.

	Brain	MTR ads
Forward process	NNO Markov chain on Q_{7P}	Demographic flow on MTR
Backward process	Postnikov tower DP	Seidel backward process
FK weight	Drug transport probability	$\omega \cdot \text{aff}$
Bracket	NNO DP (offline)	au_compiler Pass 2 (offline)
k-invariant	$k^2 = 62$	Demographic dead-zone gap
Runtime	NNO bracket lookup	Routing table lookup

Both pipelines are instances of the same pattern: the Feynman-Kac duality compiles the forward Markov process into a bracket table offline; the runtime path is a pure table lookup. The Postnikov tower records what bracket structure exists; the HMM backward process certifies that the bracket is tight under selection (genealogical convergence). The bracket at compile time is not an approximation—it is the exact Feynman-Kac expectation, evaluated once and reused indefinitely.

16. CONCLUSION

We have presented a complete architecture for lazy demographic ad placement grounded in Joyal’s Arithmetic Universe and the Fukaya A_∞ category. The three-layer design (AU attic $\rightarrow$ Hamiltonian flow $\rightarrow$ pair-of-pants Δ) reduces memory from 240 EB to ≈ 400 GB, reduces per-slot computation from $O(|S| \cdot |P|)$ to $O(D) + O(\log |P|)$, and produces flat, debuggable, SLA-compliant artifacts through an LLVM-style compilation pass.

The architecture is not merely a clever engineering optimisation. It is a fundamentally different computational model: the precomputed matrix approach requires the closed-world assumption (all combinations known before any decision is made); the AU lazy approach operates under the open-world assumption (objects materialise only when demanded by a surgery node). At planetary scale, only the latter is computationally feasible.

The pharmacodynamic correspondence (Section 10) establishes that the same mathematical structure governs both opioid transport in the mouse brain connectome and product placement in urban transit networks. The 4ti2 computation ($\beta_1 = 37$ MTR Markov circuits, 161 Graver elements) provides the algebraic certificate of the network’s toric structure, and the East Rail cross-harbour loop is confirmed as the MTR analogue of the HPF $\leftrightarrow$ sAMY dominant oscillator.

Future work: verification of $\partial^2 = 0$ for the MTR Floer complex; deployment of the hierarchical Lagrangian compiler for behavioural cohorts; extension to road networks and the full BALBc 75-node connectome.

REFERENCES

- [1] P. Seidel, *The Historical Moran Model*, arXiv:1511.05781 [math.PR], 2015.